


NUGRAPH2 PORT GUIDE

v10_06_00_01p01 → v10_06_00_04p04

ICARUS - NUGRAPH2 GNN EXTENSIONS TO 2D DECONVOLUTION DATA/MCS

S. DEY | APRIL 2026



Thank you to Giuseppe Cerati for help with debugging!

ABOUT:

This document describes the complete port of the NuGraph2 GNN extensions from icaruscode **v10_06_00_01p01** (1D deconvolution) to **v10_06_00_04p04** (2D deconvolution / WireCell). It covers all seven repositories, the specific changes required in each, common errors encountered, and their fixes.

REPOSITORIES: larrecodnn, lardataobj, larpandoracontent, larpandora, sbnanaobj, icaruscode, sbncode.

—

PART A | Higher-level overview, mapping out what to do & in what order, details of python notebooks

PART B | Full details of port

PART C | Post-port fixes (most difficult, 1D → 2D deconvolution adjustments to base code)

-

APPENDIX A | NuGraph as a part of the ICARUS/ SBN reconstruction framework

APPENDIX B | grep cheatsheet

APPENDIX C | gdb cheatsheet





A1. HIGHER-LEVEL OVERVIEW AND CONTEXT

1.1 WHAT IS BEING PORTED

The ICARUS experiment uses a Pandora-based reconstruction pipeline that runs in stages:

Stage	Description	Output
Stage 0	Raw detector data	Raw ROOT file
Stage 1	Reconstruction: clustering, Pandora, tracks, showers	Reco ROOT file
CAF	Common Analysis Format filling from Stage 1 products	Flat ROOT file

NuGraph2 is a Graph Neural Network that runs inside the Stage 1 pipeline. It assigns semantic labels (MIP, HIP, Shower, Michel, Diffuse) and a filter score to individual TPC hits, then feeds a second-pass Pandora reconstruction using only neutrino-tagged hits.

The NuGraph2 extensions were developed against `v10_06_00_01p01` (1D deconvolution). The goal of this port is to make the same NuGraph2 functionality available in `v10_06_00_04p04` (2D deconvolution / WireCell).

1.2 THE TWO-LAYER PROBLEM

Before starting any porting work, it is essential to understand that there are two independent layers of change stacked on each other:

Layer	What It Is	Key Changes
Layer 1	Base version upgrade: <code>v01p01</code> → <code>v04p04</code>	2D deconvolution, WireCell, combineHits producers, new CAF trees
Layer 2	NuGraph2 patch on top of Layer 1	New modules, FCL sequences, CAF variables, filling logic

The NuGraph patch was written assuming Layer 1 does not exist. Wherever the NuGraph code references stage1 product labels, data structures, or module names that changed in the base version upgrade, it will fail even if the patch applies cleanly. The diagnostic methodology in Section 3 characterises Layer 1 before attempting Layer 2.



A2. REPOSITORIES INVOLVED

Seven repositories carry the NuGraph2 changes. They are listed in dependency order (most upstream first). This order determines the correct port sequence.

#	Repository	Role in NuGraph Pipeline	Fork / Location	Port Order
1	larpandoracontent	Pure Pandora C++ - adds m_hitScores / m_hitScoreLabels to LArCaloHit	rtriozzi/larpandoracontent	First
2	lardataobj	LArSoft data object definitions - NuGraph class additions	leonardo-lena/lardataobj	Skip (upstream)
3	larrecodnn	Base NuGraph2 inference framework	leonardo-lena/larrecodnn	Skip (upstream)
4	larpandora	LArSoft↔Pandora interface - collects NuGraph hit labels	rtriozzi/larpandora	Second
5	icaruscode	ICARUS-specific modules, FCL sequences, stage1 integration	SBNSoftware/icaruscode	Third
6	sbnanaobj	CAF data object definitions - SRNuGraphSliceInfo, SRSlice additions	SBNSoftware/sbnanaobj	Fourth
7	sbncode	CAF filling logic - FillSliceNuGraph, FillPFPNuGraph	SBNSoftware/sbncode	Fifth



A3. DIAGNOSTIC METHODOLOGY

Run all four diagnostic steps before touching any code. These steps tell you exactly what work is needed and prevent surprises during the port.

3.1 CHECK DEPENDENCY VERSIONS UNDER BOTH BASE RELEASES

Run in a terminal (not Jupyter - env vars do not persist across Jupyter cells):

```
# v01p01 dependency versions
source /cvmfs/icarus.opensciencegrid.org/products/icarus/setup_icarus.sh
setup icaruscode v10_06_00_01p01 -q e26:prof
echo "SBNANAOBJ=$SBNANAOBJ_VERSION"
echo "SBNCODE=$SBNCODE_VERSION"
echo "LARRECODNN=$LARRECODNN_VERSION"
echo "LARPANDORA=$LARPANDORA_VERSION"
echo "LARPANDORACONTENT=$LARPANDORACONTENT_VERSION"
```

```
# v04p04 dependency versions (new terminal)
source /cvmfs/icarus.opensciencegrid.org/products/icarus/setup_icarus.sh
setup icaruscode v10_06_00_04p04 -q e26:prof
echo "SBNANAOBJ=$SBNANAOBJ_VERSION"
echo "SBNCODE=$SBNCODE_VERSION"
echo "LARRECODNN=$LARRECODNN_VERSION"
echo "LARPANDORA=$LARPANDORA_VERSION"
echo "LARPANDORACONTENT=$LARPANDORACONTENT_VERSION"
```

Reference output for this port:

Package	v01p01 Version	v04p04 Version	Changed?	Implication
sbnanaobj	v10_00_04	v10_00_05	Yes	1 patch bump - minor context drift possible
sbncode	v10_06_00_01	v10_06_00_04	Yes	3 patch bumps - manual porting may be needed
larrecodnn	v10_01_10	v10_01_10	Identical	Zero context drift - patch applies directly
larpandora	v10_00_19	v10_00_19	Identical	Zero context drift - patch applies directly
larpandoracontent	v04_15_01	v04_15_01	Identical	Zero context drift - patch applies directly

Note: Identical versions mean the patch applies against byte-for-byte identical base code.

3.2 GENERATE DIFFS FOR ALL 7 REPOS

Use the notebook [all_package_diffs_cells.ipynb](#) with version strings hardcoded (not from env vars). The notebook generates `full_diff_<repo>.patch` for each package.

Summary of diff scope (reference output for this port):



Repo	Files Changed	Patch Size	Risk	Key Notes
icaruscode	38	163.6 KB	Medium-High	Bulk is net-new NuGraph dir; existing FCL edits are the risk
sbncode	4	22.7 KB	Medium-High	CAF filling logic - API signature refactor
larpandora	6	11.0 KB	Medium	Surgical C++ changes to LArPandoraInterface/
larpandoracontent	1	9.6 KB	Low	Single header LArCaloHit.h - purely additive
sbnnanobj	5	5.6 KB	Low-Medium	New SRNuGraphSliceInfo class + SRSlice additions; watch classes_def.xml
lardataobj	0	0.0 KB	None	Already merged upstream - verify with git log
larrecodnn	0	0.0 KB	None	Already merged upstream - verify with git log

3.3 CAF VARIABLE COMPARISON (CAFVARDIFFBETWEENVERSIONS.IPYNB)

Run the notebook pointing at one vanilla v01p01 CAF file and one vanilla v04p04 CAF file. The key output is the pairwise branch presence table.

Reference output:

Comparison	Branches in A	Branches in B	Only in A	Only in B	Jaccard Similarity
v01p01 vs v04p04	2057	2106	0	49	0.977
v01p01 vs v06p03	2057	2194	0	137	0.938
v04p04 vs v06p03	2106	2194	0	88	0.960

Critical finding: $n_{\text{only_A}} = 0$ in every comparison.

No CAF branch was removed or renamed between any of these versions. Every v01p01 variable exists identically in v04p04. The 49 new branches in v04p04 are pure additions. This means the NuGraph CAF filling code in sbncode requires no variable name changes whatsoever.

3.4 STAGE I PRODUCT LABEL COMPARISON (STAGE1_CAF_PROPAGATION_NOTEBOOK.IPYNB)

Run the notebook pointing at one vanilla v01p01 and one vanilla v04p04 stage1 ROOT file. This identifies product label changes that affect the NuGraph FCL sequences.

Product Label	v01p01 NuGraph	v04p04 Vanilla	Significance
pandoraGausNuGraphRecoCryoW/E	Present	Absent	Expected - NuGraph not yet in v04p04
semantic	Present	Absent	Expected - NuGraph not yet in v04p04
filter	Present	Absent	Expected - NuGraph not yet in v04p04



ngfilteredhitsCryoE/W	Present	Absent	Expected - NuGraph not yet in v04p04
combineHitsCryoW/E	Absent	Present	NEW 2D deco step - NuGraph FCLs must use this label
400 / 600 / 300	Absent	Present	NEW WireCell tick grouping producers
cluster3DCryoW/E	Present	Present	Unchanged - safe to reference
pandoraGausCryoW/E	Present	Present	Unchanged - safe to reference
SBNShowerGausCryoW/E	Present	Present	Unchanged - safe to reference

Critical finding: combineHitsCryoW/E replaces the v01p01 1D Gauss hit source.

In v04p04, hits feeding into Pandora come from the combineHits step. Any NuGraph module or FCL that references the old hit label must be updated to use combineHitsCryoW/E. This is a targeted string substitution in `nugraph_icarus.fcl` and `stage1_icarus_defs.fcl` - not a structural refactor.



A4. RECOMMENDED PORT SEQUENCE

Follow this order strictly. Building in the wrong sequence causes linker errors because downstream packages depend on headers that have not yet been updated.

Step	Repo	Risk	Rationale
1	larpandoracontent	Low	Most upstream; identical base; single header
2	larpandora	Medium	Depends on larpandoracontent; identical base; mostly additive C++
3	sbnaobj	Low-Medium	Defines CAF types that sbncode fills - must come before sbncode
4	icaruscode	Medium-High	Stage1 FCL integration; independent of sbn* - can build in parallel
5	sbncode	Medium-High	Depends on sbnaobj for types; needs stage1 output for testing

Note: lardataobj and larrecodnn are skipped - their NuGraph changes were already merged into the base tags.



A5. TESTING CHECKPOINTS

#	Checkpoint	Command	Expected Result
1	Build success	<code>mrb b 2>&1 tail -5</code>	0 errors, 'Build complete'
2	FCL dry run	<code>lar --dry-run -c stage1_run2_icarus_NuGraphReco.fcl /dev/null</code>	Exit 0, no ERROR lines
3	Stage1 NuGraph products	<code>lar -c eventdump.fcl output.root 2>&1 grep NGMultiSlice</code>	Product entries visible
4	CAF NuGraph branches	Python: check branches with 'ng_' in name	>0 branches found
5	Score sanity	Python: <code>ng_filt_pass_frac</code> mean/min/max	Values in [0, 1]



A6. NUGRAPH2 PIPELINE DATA FLOW

This shows the full data flow within the NuGraph2-extended Stage 1 reconstruction.

Module	Role / Output
Input hits	combineHitsCryoW/E (v04p04) or gaushit (v01p01)
NCCSlicesProducer	Identifies candidate neutrino slices from input hits
NGMultiSliceInference	Runs NuGraph2 on each slice → ‘filter’ output (signal/noise score per hit) + ‘semantic’ output (MIP / HIP / Shower / Michel / Diffuse)
ngfilteredhitsProducer	Keeps only hits passing the filter score cut (configurable via NuGraphFilterCut param)
pandoraGausNuGraphReco	Second-pass Pandora reconstruction on filtered hits only
CAFMaker (sbncode)	FillSliceNuGraph() → rec.slc.ng_filt_pass_frac, rec.slc.ng_plane[3]FillPFNuGraph() → rec.slc.reco.pfp.ngscore.*



A7. QUICK REFERENCE: FEATURE BRANCH NAMES

I followed this naming convention when creating your branches so others can identify them:

Repository	Suggested Branch Name
icaruscode	feature/<username>_NuGraph2_Filter_v04p04
sbnanaobj	feature/<username>_NuGraph2_MoveVars_v04p04
sbncode	feature/<username>_CAFMaker_PandoraAfterNuGraph_v04p04
larpandora	feature/<username>_NuGraphInterface_v04p04
larpandoracontent	feature/<username>_NuGraphInterface_v04p04

This document was produced as part of the v10_06_00_01p01 → v10_06_00_04p04 NuGraph2 port effort (April 2026). The diagnostic notebooks referenced throughout should be rerun for any future port to a new base version, as the specific findings are version-specific.





B0. RECAP

0.1 WHAT CHANGED BETWEEN V01P01 AND V04P04

The key difference is the signal processing pipeline:

- ≡ v10_06_00_01p01: 1D deconvolution. NuGraph hit label = pandoraGausCryoE/W
- ≡ v10_06_00_04p04: 2D deconvolution (WireCell). NuGraph hit label = combineHitsCryoE/W

WARNING: The single most critical variable change: wherever NuGraph consumed the hit collection (HitLabel in nugraph_icarus.fcl), the label changed from pandoraGausCryoE/W to combineHitsCryoE/W. All other pandoraGausCryoE/W references (PandoraLabel for slices and PFPs) remain correct and must NOT be changed.

0.2 DEPENDENCY VERSIONS

Package	Tag	Notes
larpandoracontent	v04_15_01	Identical in both releases
larpandora	v10_00_19	Identical in both releases
sbnanaobj	v10_00_04 → v10_00_05	Changed - must use v10_00_05
sbncode	v10_06_00_01 → v10_06_00_04	Changed
icaruscode	v10_06_00_01p01 → v10_06_00_04p04	Changed

0.3 DAILY STARTUP

```
ssh -KXY icarusgpvm05.fnal.gov
sh /exp/$(id -ng)/data/users/vito/podman/start_SL7dev_jsl.sh
source /cvmfs/icarus.opensciencegrid.org/products/icarus/setup_icarus.sh
htgettoken -a htvaultprod.fnal.gov -i icarus
```

– Separate terminal –

```
ssh icarusbuild02.fnal.gov
sh /exp/$(id -ng)/data/users/vito/podman/start_SL7dev_jsl.sh
source /cvmfs/icarus.opensciencegrid.org/products/icarus/setup_icarus.sh
setup icaruscode v10_06_00_04p04 -q e26:prof
cd /exp/icarus/app/users/sdey2/dev/dev_v10_06_00_04p04_nugraph
source localProducts*/setup && mrbslp
```



Bl. LARPANDORACONTENT

Tag	v04_15_01 (identical)
Branch	feature/sdey_NuGraphInterface_v04p04
Risk	Low
Result	BUILD SUCCESS

1.1 STEPS

larpandoracontent is the same tag in both releases. The NuGraph interface files are pure additions with no conflicts.

```
❏ Check out at correct tag and create branch:
cd $MRB_SOURCE
mrb g -t v04_15_01 larpandoracontent
cd larpandoracontent
git checkout -b feature/sdey_NuGraphInterface_v04p04
❏ Apply the patch - should be clean:
git apply --reject --whitespace=fix /path/to/full_diff_larpandoracontent.patch
find . -name '*.rej' | sort
❏ Build: mrb i -j10
```

INFO: No errors encountered. Patch applied cleanly for this repository.



B2. LARPANDORA

Tag	v10_00_19 (identical)
Branch	feature/sdey_NuGraphInterface_v04p04
Risk	Low
Result	BUILD SUCCESS

Same procedure as larpandoracontent. Same tag in both releases, pure additions.

```
cd $MRB_SOURCE && mrb g -t v10_00_19 larpandora
cd larpandora
git checkout -b feature/sdey_NuGraphInterface_v04p04
git apply --reject --whitespace=fix /path/to/full_diff_larpandora.patch
mrb i -j10
```



B3. SBNANAOBJ

Tag	v10_00_05 (was v10_00_04)
Branch	feature/sdey_NuGraph2_MoveVars_v04p04
Risk	Medium
Result	BUILD SUCCESS

3.1 STEPS

- ⇒ **CRITICAL:** Check out v10_00_05, NOT v10_00_04:


```
cd $MRB_SOURCE
mrb g -t v10_00_05 sbnanaobj
cd sbnanaobj
git checkout -b feature/sdey_NuGraph2_MoveVars_v04p04
```
- ⇒ Apply the patch:


```
git apply --reject --whitespace=fix /path/to/full_diff_sbnanaobj.patch
```
- ⇒ Fix classes_def.xml checksums (critical - see Section 3.2).
- ⇒ Build.

3.2 THE CLASSES_DEF.XML CHECKSUM PROBLEM

ROOT requires exact checksums for versioned classes. The checksums from v01p01 are WRONG for v04p04 because the data model changed in v10_00_05.

WARNING: Never copy checksums from v01p01 into v04p04. They will cause ROOT schema evolution errors at runtime even if the build succeeds. Always source checksums from the v10_00_05 baseline classes_def.xml.

Correct v10_00_05 checksums:

Class	Correct Value
SRNuGraphSliceInfo	ClassVersion=10, checksum=2711257956
SRSlice	ClassVersion=23, checksum=2614728424

Key rules:

- ⇒ checksum=0 is valid as a placeholder - ROOT computes the real value on first build
- ⇒ ClassVersion in classes_def.xml must match ClassDef(ClassName, N) in the .h file
- ⇒ Preserve the full version history from v10_00_05 for SRSlice (v22 down to v1)
- ⇒ Phantom version entries from v01p01 intermediate history must be stripped

3.3 NEW FIELDS IN V10_00_05

These fields are needed by sbncode FillReco.cxx and must be present:

```
// SRSlice.h:
SRNuGraphSliceInfo ng_plane[3]; // per-plane NuGraph2 info
```




```
// SRNuGraphScore.h:
int mip_hits, hip_hits, shr_hits, mhl_hits, dif_hits;
int ng_vtx_hip_hits;
int unclustered_shr_hits;
enum NuGraphCategory { MIP, HIP, Shower, Michel, Diffuse, Unset };
```



B4. ICARUSCODE

Tag	v10_06_00_04p04
Branch	feature/sdey_NuGraph2_Filter_v04p04
Risk	High
Result	BUILD SUCCESS

4.1 OVERVIEW

icaruscode requires the most manual work. The patch from v01p01 will have many rejects because FCL files changed significantly between releases. Always inspect the actual v04p04 file content before making edits.

4.2 CHECKOUT

```
cd $MRB_SOURCE
mrb g -t v10_06_00_04p04 icaruscode
cd icaruscode
git checkout -b feature/sdey_NuGraph2_Filter_v04p04
```

4.3 APPLY PATCH THEN RESET MODIFIED FILES

Apply with `--reject` to extract the pure additions (NuGraph/ directory, new FCLs), then reset the partially-modified existing files to a clean baseline:

```
git apply --reject --whitespace=fix /path/to/full_diff_icaruscode.patch
git checkout -- CMakeLists.txt fcl/caf/cafmaker_defs.fcl \
  fcl/reco/Definitions/stagel_icarus_defs.fcl \
  fcl/reco/Stagel/data/stagel_run2_1d_icarus.fcl \
  fcl/reco/Stagel/data/stagel_run2_icarus.fcl \
  fcl/reco/Stagel/mc/stagel_run2_1d_icarus_MC.fcl \
  fcl/reco/Stagel/mc/stagel_run2_icarus_MC.fcl \
  icaruscode/TPC/CMakeLists.txt \
  icaruscode/TPC/Tracking/ICARUSPandora/pandoramodules_icarus.fcl
find . -name '*.rej' -delete
```

INFO: The new untracked files (NuGraph/ directory, NuGraphReco FCLs) are correct and should be kept as-is.

4.4 MANUAL EDITS - 9 FILES

File 1: CMakeLists.txt (top-level)

Add `larrecodnn` after `larpandoracontent`. Check current line number first:

```
grep -n 'find_package' CMakeLists.txt | grep -A2 'larpandoracontent'
# Insert after larpandoracontent line (typically 74):
sed -i '74a find_package(larrecodnn REQUIRED )' CMakeLists.txt
grep -n 'larrecodnn' CMakeLists.txt # verify exactly once
```

File 2: icaruscode/TPC/CMakeLists.txt



Add NuGraph subdirectory after SPAna:

```
grep -n 'add_subdirectory' icaruscode/TPC/CMakeLists.txt
# Add after SPAna (typically line 11):
sed -i '11a add_subdirectory(NuGraph)' icaruscode/TPC/CMakeLists.txt
grep -n 'NuGraph' icaruscode/TPC/CMakeLists.txt # must be exactly once!
```

WARNING: CRITICAL: If NuGraph appears twice in this file you will get a CMake error 'binary directory already used'. Check with grep first and ensure NuGraph appears EXACTLY ONCE.

File 3: pandoramodules_icarus.fcl

Add 2 commented NuGraph lines after the closing brace of icarus_basicpandora (check line ~23):

```
sed -n '20,28p' icaruscode/TPC/Tracking/ICARUSPandora/pandoramodules_icarus.fcl
# Manually add after closing } of icarus_basicpandora:
# NuGraph settings (enable when running NuGraph):
# icarus_basicpandora.ShouldCollectHitPredictions: true
```

File 4: fcl/caf/cafmaker_defs.fcl

Add all NuGraph CAF parameters:

```
grep -n 'TrackHitFillRREndCut' fcl/caf/cafmaker_defs.fcl # find line number
# Insert after that line (typically 318):
sed -i '318a cafmaker.NuGraphSlicesLabel: "NCCSlicesCryoE"\
ncafmaker.NuGraphFilterLabel: "filter"\
ncafmaker.NuGraphSemanticLabel: "semantic"\
ncafmaker.UsePandoraAfterNuGraph: true\
ncafmaker.NuGraphFilterCut: 0.5\
ncafmaker.NuGraphHIPTagWireDist: 10\
ncafmaker.NuGraphHIPTagTickDist: 50' fcl/caf/cafmaker_defs.fcl
grep -n 'NuGraph' fcl/caf/cafmaker_defs.fcl # verify 7 lines
```

File 5: stage1_icarus_defs.fcl (the big one)

Requires 3 groups of additions. Use a Python script for reliability. Add:

⇒ Two #include lines after tpcpmtbarycentermatch_config.fcl:

```
#include "pandoramodules_icarus.fcl"
#include "nugraph_icarus.fcl"
```

⇒ NuGraph producers in icarus_stage1_producers block (after tpcpmtbarycentermatchCryoW):

```
## NuGraph producers
nuslhitsCryoE/W, NuGraphCryoE/W, NCCSlicesCryoE/W,
NGMultiSliceCryoE/W, ngfilteredhitsCryoE/W,
pandoraGausNuGraphRecoCryoE/W: { @table::icarus_pandora
  HitFinderModuleLabel: ngfilteredhitsCryoE/W }
```

⇒ postPandora sequences after icarus_pandora_Gauss:

```
icarus_reco_postPandoraGausCryoW/E: [ NCCSlices, NGMultiSlice ]
icarus_reco_postPandoraGausCryoW/E_NuGraphReco: [ NCCSlices, NGMultiSlice,
  ngfilteredhits, pandoraGausNuGraphReco, pandoraTrack, pandoraKalman, SBNShower ]
```



```
icarus_postPandora_Gauss: [ postPandoraGausCryoE, postPandoraGausCryoW ]
icarus_postPandora_Gauss_NuGraphReco: [ ...NuGraphReco variants ]
```

Files 6-9: The 4 Stage1 Run FCLs

Insert @sequence::icarus_postPandora_Gauss after pandora/Gauss sequences, BEFORE icarus_reco_fm:

File	Insertion Point
stage1_run2_icarus.fcl	After @sequence::icarus_pandora_Gauss
stage1_run2_1d_icarus.fcl	After @sequence::icarus_pandora_Gauss
stage1_run2_icarus_MC.fcl	After @sequence::icarus_reco_Gauss2D_CryoW
stage1_run2_1d_icarus_MC.fcl	After @sequence::icarus_reco_Gauss1D_CryoW

WARNING: For MC FCLs, postPandora_Gauss goes AFTER both CryoE and CryoW sequences. Final order: Gauss_CryoE, Gauss_CryoW, postPandora_Gauss, icarus_reco_fm. Never between CryoE and CryoW.

```
# Verify all 4 FCLs after editing:
grep -n 'postPandora_Gauss\|pandora_Gauss\|Gauss2D_CryoW\|reco_fm' \
  fcl/reco/Stage1/data/stage1_run2_icarus.fcl \
  fcl/reco/Stage1/data/stage1_run2_1d_icarus.fcl \
  fcl/reco/Stage1/mc/stage1_run2_icarus_MC.fcl \
  fcl/reco/Stage1/mc/stage1_run2_1d_icarus_MC.fcl
```

4.5 THE CRITICAL V04P04 VARIABLE FIX

After applying the patch, nugraph_icarus.fcl has the old v01p01 hit label. Fix it:

```
grep -n 'HitLabel' icaruscode/TPC/NuGraph/nugraph_icarus.fcl
# Fix both labels:
sed -i 's/HitLabel: "pandoraGausCryoE"/HitLabel: "combineHitsCryoE"/' \
  icaruscode/TPC/NuGraph/nugraph_icarus.fcl
sed -i 's/HitLabel: "pandoraGausCryoW"/HitLabel: "combineHitsCryoW"/' \
  icaruscode/TPC/NuGraph/nugraph_icarus.fcl
# Verify:
grep -n 'HitLabel' icaruscode/TPC/NuGraph/nugraph_icarus.fcl
```

NOTE: PandoraLabel references (pandoraGausCryoE/W) in nugraph_icarus.fcl and C++ module defaults are CORRECT - they refer to Pandora reco products (slices/PFPs), not the hit collection. Do NOT change them.

4.6 COMMIT AND BUILD

```
git add -A
git commit -m 'Port NuGraph2 icaruscode changes from v01p01 to v04p04'
cd $MRB_BUILDDIR && mrb i -j10 2>&1 | tee build_icaruscode.log
```

4.7 COMMON ERRORS



Error	Fix
CMake: binary directory already used	Duplicate add_subdirectory(NuGraph). Remove duplicate, delete build/icaruscode/, rerun mrb i.
CMake: subdirectory inclusion order	Run mrb uc from \$MRB_SOURCE.
ROOT schema evolution error	Wrong checksums in classes_def.xml. Use v10_00_05 checksums only.
FCL: unknown sequence icarus_postPandora_Gauss	Check stage1_icarus_defs.fcl includes both pandoramodules_icarus.fcl and nugraph_icarus.fcl.
FCL: unknown label combineHits*	Check stage1_icarus_defs.fcl has combineHitsCryoE/W producers defined.



B5. SBNCODE

Tag	v10_06_00_04
Branch	feature/sdey_CAFMaker_PandoraAfterNuGraph_v04p04
Risk	Medium
Result	BUILD SUCCESS

5.1 OVERVIEW

sbncode changes are in 4 files in sbncode/CAFMaker/. The v04p04 patch significantly refactors NuGraph data retrieval: switching from a pre-fetched slice hit map to per-slice art::FindOneP lookups, and adding plane-by-plane semantic variables.

File	Changes
CAFMakerParams.h	Rename NuGraphSliceHitLabel -> NuGraphSlicesLabel. Add 4 new params.
CAFMaker_module.cc	Remove old slice hit map handles. Add nuGraphSlices vector. Replace FillSlice/PFPNuGraph calls with art::FindOneP lookups.
FillReco.cxx	Remove ng_filter_cut global. Rewrite FillSliceNuGraph with plane-by-plane vars. Refactor FillPFPNuGraph.
FillReco.h	Update function signatures.

5.2 CHECKOUT AND APPLY

```
cd $MRB_SOURCE
mrb g -t v10_06_00_04 sbncode
cd sbncode
git checkout -b feature/sdey_CAFMaker_PandoraAfterNuGraph_v04p04
git apply --reject --whitespace=fix /path/to/full_diff_sbncode.patch
find . -name '*.rej' | sort # should be empty
```

SUCCESS: In the v04p04 port, all 4 sbncode files applied cleanly with zero rejects.

5.3 KEY API CHANGES (REFERENCE IF MANUAL EDITING NEEDED)

5.3.1 CAFMakerParams.h - Renamed and New Params

```
// OLD: NuGraphSliceHitLabel
// NEW: NuGraphSlicesLabel
// NEW ADDITIONS:
Atom<bool> UsePandoraAfterNuGraph { Name(...), false }
```



```
Atom<float> NuGraphFilterCut      { Name(...), 0.5  }
Atom<float> NuGraphHIPTagWireDist { Name(...), 10   }
Atom<float> NuGraphHIPTagTickDist { Name(...), 50   }
```

5.3.2 FillSliceNuGraph - New Signature

```
// OLD: takes sliceHitsMap, uses global ng_filter_cut
void FillSliceNuGraph(inputHits, sliceHitsMap, ngFilter, ngSemantic, slice);

// NEW: loses sliceHitsMap, gains fmPFPPartHits + vertex + filter_cut
void FillSliceNuGraph(inputHits, ngFilter, ngSemantic,
    fmPFPPartHits, vtx_wire[3], vtx_tick[3],
    vtx_wire_dist, vtx_tick_dist, filter_cut, slice);
```

5.3.3 FillPFPNuGraph - New Signature

```
// OLD: takes sliceHitsMap + separate pfpHits
void FillPFPNuGraph(sliceHitsMap, ngFilter, ngSemantic, pfpHits, srpfp);

// NEW: takes pfpHits directly, adds filter_cut param
void FillPFPNuGraph(pfpHits, ngFilter, ngSemantic, filter_cut, srpfp);
```

5.3.4 Hit Lookup: std::find -> art::FindOneP

```
// OLD: global prefetch + std::find
GetByLabelIfExists(evt, NuGraphSliceHitLabel, handle);
auto it = std::find(sliceHitsMap.begin(), sliceHitsMap.end(), hit.key());

// NEW: per-slice art::FindOneP
art::FindOneP<anab::FeatureVector<1>> findOneFilter(
    slcHits, evt, NuGraphFilterLabel + slice_tag_suff + ":" + instance);
for (size_t i = 0; i < slcHits.size(); ++i)
    ng2_filter_vec.emplace_back(findOneFilter.at(i));
```

5.3.5 nuGraphSlices Logic

```
std::vector<art::Ptr<recob::Slice>> nuGraphSlices;
if (fParams.UsePandoraAfterNuGraph())
    nuGraphSlices = slices; // all slices
else
    nuGraphSlices = evt.getProduct<...>(NuGraphSlicesLabel + pandora_tag_suffix);

// Then guard all NuGraph filling:
if (std::find(nuGraphSlices.begin(), nuGraphSlices.end(), slice) !=
    nuGraphSlices.end()) {
    // fill NuGraph variables for this slice
}
```

5.4 UPDATE CAFMAKER_DEFS.FCL IN ICARUSCODE

After patching sbncode, update icaruscode cafmaker_defs.fcl with all NuGraph parameters:

```
# In fcl/caf/cafmaker_defs.fcl - full NuGraph block:
```



```
cafmaker.NuGraphSlicesLabel:      "NCCSlicesCryoE"
cafmaker.NuGraphFilterLabel:      "filter"
cafmaker.NuGraphSemanticLabel:    "semantic"
cafmaker.UsePandoraAfterNuGraph:  true
cafmaker.NuGraphFilterCut:        0.5
cafmaker.NuGraphHIPTagWireDist:   10
cafmaker.NuGraphHIPTagTickDist:   50
```

WARNING: NuGraphSliceHitLabel is gone. Using the old name will cause a fatal FHiCL validation error at runtime.

5.5 COMMIT AND BUILD

```
cd $MRB_SOURCE/sbncode
git add sbncode/CAFMaker/CAFMakerParams.h sbncode/CAFMaker/CAFMaker_module.cc \
      sbncode/CAFMaker/FillReco.cxx sbncode/CAFMaker/FillReco.h
git commit -m 'Port NuGraph2 CAFMaker changes from v01p01 to v04p04'

# Also amend icaruscode to include cafmaker_defs.fcl:
cd $MRB_SOURCE/icaruscode
git add fcl/caf/cafmaker_defs.fcl && git commit --amend --no-edit

# Run mrb uc to get correct build order, then build:
cd $MRB_SOURCE && mrb uc
mrbsetenv
cd $MRB_BUILDDIR && mrb i -j10 2>&1 | tee build_sbncode.log
```

5.6 COMMON ERRORS

Error	Fix
CMake subdirectory order error	Run mrb uc after adding sbncode. Correct order: larpandoracontent, larpandora, sbnanaobj, icaruscode, sbncode - just manually change this since mrb uc doesn't always do this correctly.
FHiCL: unknown key NuGraphSliceHitLabel	Update cafmaker_defs.fcl to NuGraphSlicesLabel.
Compile: ng_filter_cut undeclared	Remove the old global from FillReco.cxx - it is now a function parameter.
art exception on FindOneP	Check NuGraphFilterLabel and NuGraphSemanticLabel are set correctly in FCL.



B6. FINAL CHECKLIST

6.1 BUILD ORDER

Order	Command
1. larpandoracontent	<code>mrbs g -t v04_15_01 larpandoracontent</code>
2. larpandora	<code>mrbs g -t v10_00_19 larpandora</code>
3. sbnanaobj	<code>mrbs g -t v10_00_05 sbnanaobj (NOT v10_00_04!)</code>
4. icaruscode	<code>mrbs g -t v10_06_00_04p04 icaruscode</code>
5. sbncode	<code>mrbs g -t v10_06_00_04 sbncode</code>

6.2 PRE-BUILD CHECKLIST

- ⊟ icaruscode/TPC/CMakeLists.txt has `add_subdirectory(NuGraph)` exactly ONCE
- ⊟ `nugraph_icarus.fcl:HitLabel = combineHitsCryoE/W` (not `pandoraGausCryoE/W`)
- ⊟ `cafmaker_defs.fcl`: uses `NuGraphSlicesLabel` (not `NuGraphSliceHitLabel`)
- ⊟ `cafmaker_defs.fcl`: has all 7 `NuGraph` parameters
- ⊟ `stage1_icarus_defs.fcl`: includes `pandoramodules_icarus.fcl` AND `nugraph_icarus.fcl`
- ⊟ All 4 stage1 run FCLs: have `@sequence::icarus_postPandora_Gauss` in correct position
- ⊟ `sbnanaobj_classes_def.xml`: uses `v10_00_05` checksums
- ⊟ `mrbs uc` run after adding all 5 packages (though typically it still gives you a messed up order so put it in manually or just follow my order)

6.3 VARIABLE CHANGE SUMMARY

Item	v01p01	v04p04
<code>nugraph_icarus.fcl</code> <code>HitLabel</code>	<code>pandoraGausCryoE/W</code>	<code>combineHitsCryoE/W</code>
CAFMakerParams param name	<code>NuGraphSliceHitLabel</code>	<code>NuGraphSlicesLabel</code>
New param	(absent)	<code>UsePandoraAfterNuGraph = true</code>
New param	(absent)	<code>NuGraphFilterCut = 0.5</code>
New param	(absent)	<code>NuGraphHIPTagWireDist = 10</code>
New param	(absent)	<code>NuGraphHIPTagTickDist = 50</code>
FillReco.cxx global	<code>const float ng_filter_cut = 0.5</code>	Removed - function param
sbnanaobj tag	<code>v10_00_04</code>	<code>v10_00_05</code>



6.4 BRANCH NAMES

Repository	Branch
larpandoracontent	feature/sdey_NuGraphInterface_v04p04
larpandora	feature/sdey_NuGraphInterface_v04p04
sbnanaobj	feature/sdey_NuGraph2_MoveVars_v04p04
icaruscode	feature/sdey_NuGraph2_Filter_v04p04
sbncode	feature/sdey_CAFMaker_PandoraAfterNuGraph_v04p04

SUCCESS: All 5 repositories successfully ported and built for v10_06_00_04p04. Next step: run a test job with `stage1_run2_icarus_NuGraphReco.fcl` to verify runtime correctness of the full NuGraph2 pipeline.





C. POST-BUILD RUNTIME DEBUGGING AND FINAL FIXES

C0. CONTEXT & STARTING POINT

This section covers the runtime debugging phase that followed the successful build of all five ported repositories. At the start of Part C, the build completes without errors and art can load the FCL. The issues from this point forward are runtime crashes and incorrect data product associations, only visible when lar is actually executed.

Parts A and B of this guide covered the port itself: which files to change, in which repos, and in what order. **Part C** covers what happened next, how it was diagnosed, and exactly what was fixed.



CI. FIRST RUNTIME TEST — SEGFAULT (HEAP CORRUPTION)

CI.1 THE RUN

After a successful build, the standard run command was:

```
lar -c stage1_run2_icarus_overlay_NuGraphReco.fcl -n 1 \
-s /pnfs/sbn/data/sbn_fd/poms_production/mc/2025A_ICARUS_Overlays_BNB_RUN2/\
September/v10_06_00_04p04/stage0/00/02/\
stage0_run2_wcdnn_icarus_overlay_26910969_7314-...root \
--process-name redo -o stage1-nugraph-test.root \
--trace > stage1_nugraph.log 2>&1
```

The `--trace` flag prints each module as it starts and finishes — the fastest way to find where a crash occurs.

The job crashed with a segfault. The `--trace` output showed:

```
+++++++ module for event: NGMultiSliceCryoE
+++++++ finished for event: NGMultiSliceCryoE    <- NuGraph inference SUCCEEDS
+++++++ module for event: ngfilteredhitsCryoE    <- log cuts off here, core
dump
```

The job died at the start of `ngfilteredhitsCryoE` (ICARUSFilteredNuSliceHitsProducer). NuGraph inference itself completed successfully.

CI.2 DIAGNOSTIC APPROACH — THREE TOOLS

Three complementary tools were used to diagnose the crash. Each reveals a different layer of the problem.

Tool 1: GDB — Where Did It Crash?

GDB catches the crash and prints a backtrace showing the call stack at the moment of the segfault. Before running, unset Python environment variables that cause GDB's embedded Python to fail:

```
unset PYTHONPATH PYTHONHOME

gdb -ex "set logging file gdb_nugraph_bt.log" \
    -ex "set logging on" \
    -ex "run" \
    -ex "bt" \
    -ex "bt full" \
    -ex "quit" \
    --args lar -c stage1_run2_icarus_overlay_NuGraphReco.fcl -n 1 \
    -s <input_file> --process-name redo -o /dev/null 2>&1
```

Key GDB backtrace frames:

```
#0  _int_malloc                                <- heap allocator corrupted
#1  malloc
#2  operator new
#3  art::detail::IPRHelper<recob::Slice, recob::Hit...>
```



```
<- FindManyP<Hit> construction inside ICARUSNuGraphInference_module.so
#4 ICARUSFilteredNuSliceHitsProducer::produce
```

The crash is in `_int_malloc` — the heap allocator itself. This is a classic heap corruption pattern: the crash site (malloc) and the bug site (the bad write) are different locations. The actual out-of-bounds write happened upstream, and `ICARUSFilteredNuSliceHitsProducer` is simply the first module to call malloc after the corruption.

Tool 2: Valgrind — Where Was the Bad Write?

Valgrind intercepts every memory read and write at the CPU level, catching the invalid write at the exact moment it happens. The tradeoff is **10-100x** slower execution. Always run in the background with `&`.

```
unset PYTHONPATH PYTHONHOME
setup valgrind v3_12_0 -q "Linux64bit+3.10-2.17"

valgrind --tool=memcheck \
  --error-limit=no \
  --track-origins=yes \
  --num-callers=30 \
  lar -c stage1_run2_icarus_overlay_NuGraphReco.fcl -n 1 \
  -s <input_file> --process-name redo -o /dev/null \
  > valgrind_nugraph.log 2>&1 &

# Monitor progress:
tail -f valgrind_nugraph.log
```

Flag	Meaning
<code>--tool=memcheck</code>	Use the memory checking tool (default but explicit)
<code>--error-limit=no</code>	Do not stop reporting after 1000 errors
<code>--track-origins=yes</code>	Report where uninitialised values came from (slower but more informative)
<code>--num-callers=30</code>	Show 30 stack frames per error instead of default 12
<code>&</code>	Run in background, freeing the terminal

Valgrind confirmed three invalid reads inside `HitMerger` (`makeWireAssns`, `makeRawDigitAssns`, `makeChanROIAssns`) and zero invalid reads inside any NuGraph code. The `HitMerger` errors were identical in both the NuGraph run and a vanilla v04p04 run - confirming a pre-existing bug in base icaruscode, not something introduced by the port.

Tool 3: Debug Build — For Future Reference

GDB output from a prof (optimised) build often shows variables as `<optimised out>`. A debug build compiles with `-O0 -g`, giving full symbol visibility. To add a debug build alongside the existing prof build:

```
# From the top-level dev area directory:
cd /path/to/dev_v10_06_00_04p04_nugraph

# -T names the new build directory, -f reuses existing srcs/
# mrb newDev -T debug_build -q e26:debug -f

# Source the debug environment and build:
```



```
source localProducts*debug_build*/setup
mrbsetenv
mrb i -j8 2>&1 | tee build_debug.log
```

WARNING: A debug build is significantly larger and slower than a prof build. Only use it when you need to inspect local variable values in GDB. For this port, the prof build combined with valgrind and printouts was sufficient.



C2. FIXES APPLIED

C2.1 FIX: HITMERGER VECTOR REALLOCATION BUG

Note: This bug exists in vanilla v04p04 and is not introduced by the NuGraph port. It should be submitted as a PR to base icaruscode.

`HitMerger` is a v04p04-specific module (absent in v01p01) that merges per-TPC hit collections from WireCell 2D deconvolution into a single collection per cryostat.

```
gaushit2dTPCEW  -->|
                  |--> HitMerger --> combineHitsCryoE
gaushit2dTPCEE  -->|
```

File: icaruscode/TPC/SignalProcessing/HitFinder/HitMerger_module.cc

Root cause: `HitMerger::produce()` stores raw `recob::Hit*` pointers as keys in `recobHitToPtrMap`. With two input labels (`gaushit2dTPCEW` and `gaushit2dTPCEE`), `emplace_back` on the second pass triggers reallocation of `outputHitPtrVec`, moving the vector to a new memory location. All raw pointers stored from the first pass are now dangling. The `makeWireAssns`, `makeRawDigitAssns`, and `makeChanROIAssns` functions then dereference those dangling pointers - reading from freed memory.

Fix: pre-count all hits and call `reserve()` before the loop so the vector never reallocates.

```
// Sparshita Dey 2026 -- pre-count to prevent reallocation invalidating raw
pointers
size_t totalHits = 0;
for (const auto& inputTag : HitMergerfHitProducerLabelVec) {
    art::Handle<std::vector<recob::Hit>> hitHandle;
    evt.getByLabel(inputTag, hitHandle);
    if (hitHandle.isValid()) totalHits += hitHandle->size();
}
outputHitPtrVec->reserve(totalHits);
```

After applying this fix and rebuilding, valgrind confirmed the `HitMerger` errors were gone.

C2.2 FIX: NGMULTISLICECRYOW WRONG HITINPUT (FCL BUG)

File: icaruscode/TPC/NuGraph/nugraph_icarus.fcl

`NGMultiSliceCryoW` was defined as `@local::NGMultiSliceCryoE` and inherited `hitInput: nuslhitsCryoE` for its loader and decoder tools - feeding the west cryostat NuGraph with east cryostat hits.

Fix: add explicit overrides after the `NGMultiSliceCryoW` block in all three FCL copies (`srcs/`, `localProducts*/`, `build_slf7*/`):

```
NGMultiSliceCryoW.LoaderTool.hitInput: "nuslhitsCryoW"
NGMultiSliceCryoW.LoaderTool.spsInput: "nuslhitsCryoW"
```




WARNING: Remember: `lar` reads from `localProducts*/`, not `srcs/`. All three copies of every FCL must be edited for the change to take effect at runtime.

C2.3 FIX: SWITCH PANDORA TO CLUSTER3DCRYOE (ARCHITECTURE FIX)

C2.3a Why This Was Needed

In the initial port, Pandora read hits from `combineHitsCryoE` (produced by `HitMerger` from `gaushit2dTPCEW/EE`). NuGraph's decoder tools, however, read `SpacePoints` from `cluster3DCryoE` (produced from `gaushitPT2dTPCEW/EE` - the pass-through hits). These are entirely separate hit collections with different `art::Ptr` keys.

The consequence: when NuGraph's `FilterDecoder::writeToEvent` used `idsmap` entries (`art::Ptr` keys into `cluster3DCryoE`) to index into `filtcol` (sized by the count of hits from `combineHitsCryoE`), it wrote out of bounds. This is the heap corruption that GDB caught in `_int_malloc`.

Giuseppe confirmed the correct fix: NuGraph requires a single consistent hit source feeding Pandora, the loader, and the decoder. The right choice is `cluster3DCryoE`, which is the source the NuGraph decoder tools already expected.

C2.3b The Fix

File: `icaruscode/fcl/reco/Definitions/stagel_icarus_defs.fcl`, lines 426 and 434

```
# Change from:
icarus_stagel_producers.pandoraGausCryoW.HitFinderModuleLabel:
"combineHitsCryoW"
icarus_stagel_producers.pandoraGausCryoE.HitFinderModuleLabel:
"combineHitsCryoE"

# Change to:
icarus_stagel_producers.pandoraGausCryoW.HitFinderModuleLabel: "cluster3DCryoW"
icarus_stagel_producers.pandoraGausCryoE.HitFinderModuleLabel: "cluster3DCryoE"
```

Apply using `sed` across all three FCL copies:

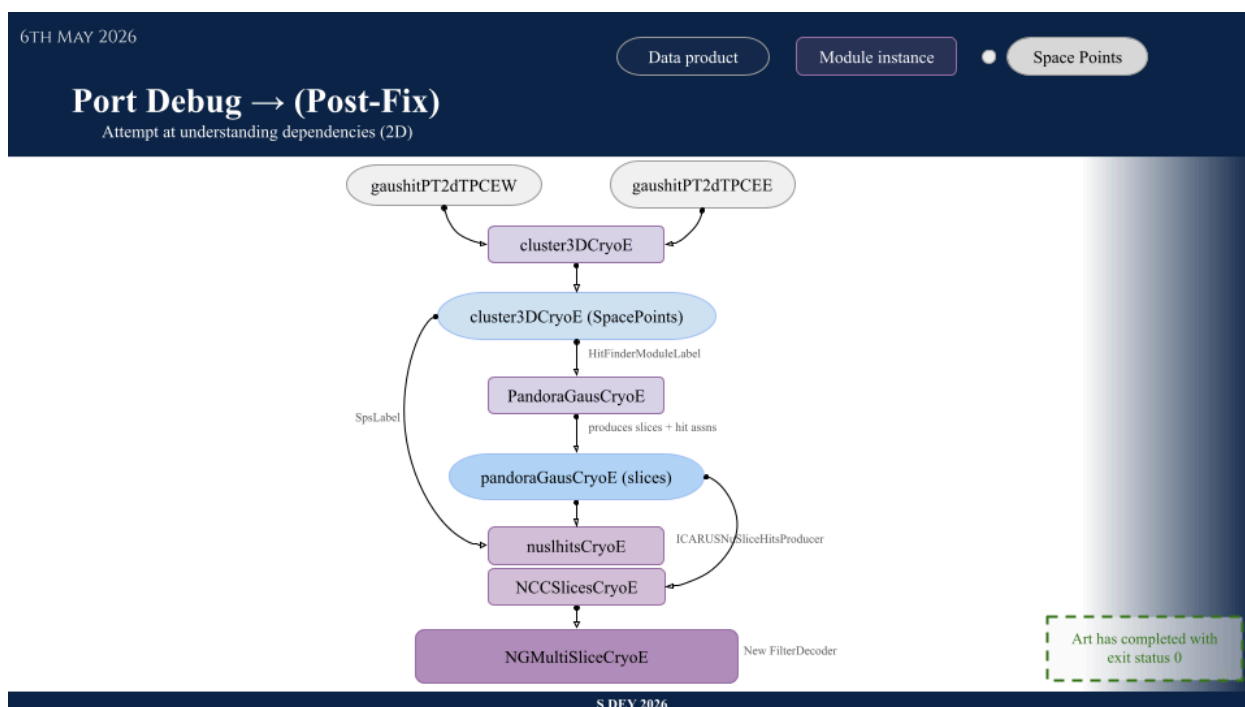
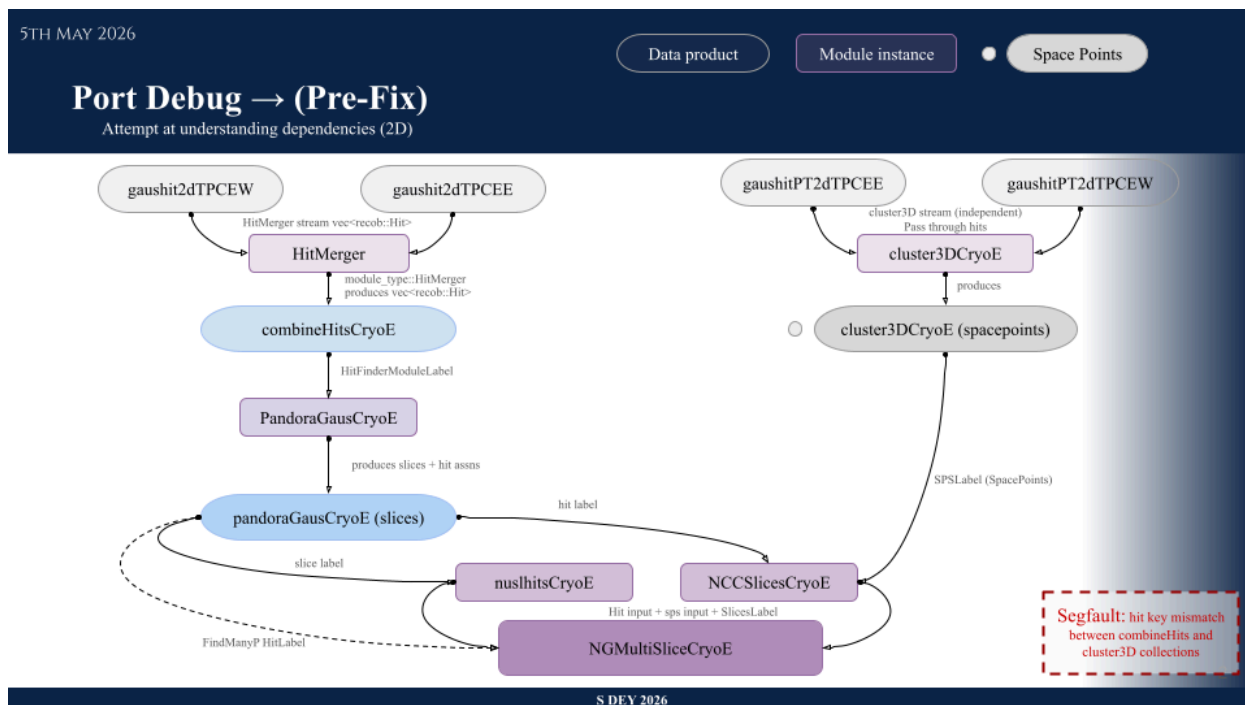
```
for fcl in \
  srcs/icaruscode/fcl/reco/Definitions/stagel_icarus_defs.fcl \
  localProducts*/icaruscode/*/job/stagel_icarus_defs.fcl \
  build_slf7.x86_64/icaruscode/job/stagel_icarus_defs.fcl; do
  sed -i
  's/pandoraGausCryoW.HitFinderModuleLabel:.*"combineHitsCryoW"/pandoraGausCryoW.
HitFinderModuleLabel: "cluster3DCryoW"/' $fcl
  sed -i
  's/pandoraGausCryoE.HitFinderModuleLabel:.*"combineHitsCryoE"/pandoraGausCryoE.
HitFinderModuleLabel: "cluster3DCryoE"/' $fcl
done

# Verify (should show cluster3DCryoE/W in all 6 lines):
grep -n 'HitFinderModuleLabel' \
  srcs/icaruscode/fcl/reco/Definitions/stagel_icarus_defs.fcl \
  localProducts*/icaruscode/*/job/stagel_icarus_defs.fcl \
  build_slf7.x86_64/icaruscode/job/stagel_icarus_defs.fcl
```

Note: `HitMerger` still runs and produces `combineHitsCryoE/W`, but these collections are no longer consumed by Pandora or NuGraph. They are unused downstream.



C2.3c Pipeline: Before and After



Note: See accompanying diagrams showing the before/after pipeline data flow.



C2.4 FIX: LARDATAOBJ AND LARRECODNN FROM LEONARDO'S BRANCH

C2.4a The FilterDecoder Bug

Even with the `cluster3D` switch, the pipeline crashed in `FilterDecoder::writeToEvent`. The version of `FilterDecoder_tool.cc` on `cvmfs (larrecodnn v10_01_10)` had an indexing bug:

```
// filtcol is sized by COUNT of hits in slices (e.g. 200):
size_t size = 0;
for (auto& v : idsmmap) size += v.size();
auto filtcol = make_unique<vector<FeatureVector<1>>>(size, ...);

// But idx is a raw art::Ptr KEY into the full hit collection (e.g. 4532):
size_t idx = idsmmap[p][i];
(*filtcol)[idx] = FeatureVector<1>(input); // CRASH if 4532 > 200
```

The raw `art::Ptr` key (which indexes into the full hit collection) was used as a direct positional index into `filtcol` (sized only by the number of hits in slices). If any hit has a key larger than the count of processed hits, this writes out of bounds.

C2.4b The Fix — Leonardo's Branch

Leonardo Lena's `feature/nugraph-multislice` branch of `larrecodnn` contains a corrected `FilterDecoder_tool.cc` that builds a sorted key list and uses `std::find + std::distance` to convert raw keys to correct positional indices. It also introduces `art::Assns<FeatureVector<1>, recob::Hit>` for proper hit-score associations. Check out both `lardataobj` and `larrecodnn` from Leonardo's branch:

```
cd $MRB_SOURCE

# lardataobj
mrbs g -t v10_01_00 lardataobj
cd lardataobj
git remote add llena https://github.com/leonardo-lena/lardataobj.git
git fetch llena
git checkout -B feature-nugraph-multislice llena/feature/nugraph-multislice
cd ..

# larrecodnn
mrbs g -t v10_01_10 larrecodnn
cd larrecodnn
git remote add llena https://github.com/leonardo-lena/larrecodnn.git
git fetch llena
git checkout -B feature-nugraph-multislice llena/feature/nugraph-multislice
cd ..

# Rebuild:
cd $MRB_BUILDDIR
mrbs setenv
mrbs i -j8 2>&1 | tee build_all.log
```

WARNING: `mrbs uc` does not reliably produce the correct build order when `lardataobj` and `larrecodnn` are added (to fix this would mean editing up dependencies for all packages → not worth it! Edit `CMakeLists.txt` manually. Correct order: `larpandoracontent`, `lardataobj`, `larrecodnn`, `larpandora`, `sbnanaobj`, `icaruscode`, `sbncode`. The build will either stall or you will get dependency errors if the order is wrong - if this happens just correct it according to the error message, then `mrbs setenv`, then rebuild (`mrbs i -j20`).



C3. FIRST SUCCESSFUL END-TO-END RUN

SUCCESS: Art has completed and will exit with status 0. Events total = 1 passed = 1 failed = 0.

After all four fixes were applied and the build completed, the full NuGraph2 pipeline ran successfully for the first time on v04p04. All modules completed with 0 errors:

```
TrigReport Events total = 1 passed = 1 failed = 0
TimeReport CPU = 34.797811 Real = 44.653458
MemReport VmPeak = 4620.84 VmHWM = 2660.17
Art has completed and will exit with status 0.
```



C4. NUGRAPH2 HIT FILTER STATISTICS

Run on 5 MC overlay events (BNB RUN2, v10_06_00_04p04 Stage 0 input). Filter score cut = 0.5. Hit counts reported by ICARUSFilteredNuSliceHitsProducer.

```
grep "Number of hits" stage1_nugraph_5events.log
```

Event	Cryostat	Hits before	1D model after	1D % kept	2D model after	2D % kept
1	CryoE	2165	1624	75.0%	1716	79.3%
1	CryoW	2961	2376	80.2%	2522	85.2%
2	CryoE	2793	2327	83.3%	2458	88.0%
2	CryoW	3629	3073	84.7%	2168	59.7%
3	CryoE	4804	3773	78.5%	3858	80.3%
3	CryoW	3989	3368	84.4%	3435	86.1%
4	CryoE	860	487	56.6%	554	64.4%
4	CryoW	505	224	44.4%	352	69.7%
5	CryoE	2983	1940	65.0%	2530	84.8%
5	CryoW	2568	2030	79.0%	2177	84.8%
Overall		27257	21222	77.9%	21770	79.9%

Note: Event 2 CryoW is an outlier at 59.7% kept with the 2D model (vs 84.7% with the 1D model). This warrants further study on more events.

Note: The 1D model column is shown for reference only. It should not be used with 2D deconvolution inputs. See Section C5.



C5. NUGRAPH2 MODEL

NuGraph2 uses a PyTorch model file (.pt) for inference, loaded at runtime via `FW_SEARCH_PATH`. No rebuild is needed to switch models. The `modelFileName` parameter in `nugraph_icarus.fcl` controls which file is loaded.

	v10_06_00_01p01 (1D)	v10_06_00_04p04 (2D)
Model file	model_mpvmpmr_bnb_numu_cos.pt	model_mpvmpmr_1201.pt
Location	cvmfs icarus_data/NuGraph/	Local (not yet on cvmfs)
Training sample	BNB numu + cosmics (1D hits)	MPVMMPR (2D WireCell hits)
Semantic labels	MIP, HIP, Shower, Michel, Diffuse	MIP, HIP, Shower, Michel, Diffuse

To use the 2D model, add its directory to `FW_SEARCH_PATH` before running:

```
export FW_SEARCH_PATH=/path/to/model/directory:$FW_SEARCH_PATH
```

WARNING: The 2D model (model_mpvmpmr_1201.pt) is not yet on `cvmfs`. It must be present in a directory on `FW_SEARCH_PATH` at runtime. This should be submitted to `icarus_data` for production use.



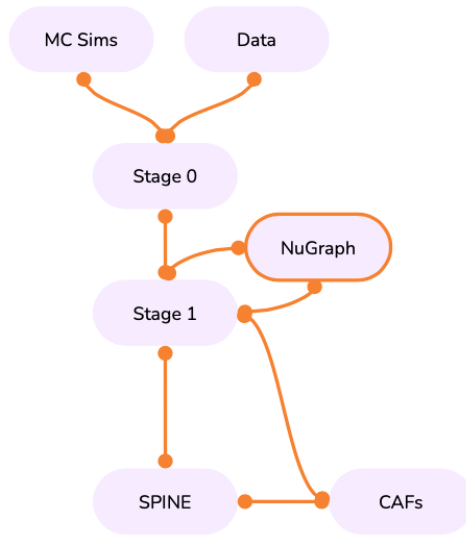
C6. CHANGE SUMMARY

Fix	File	Type	Notes
HitMerger reserve()	HitMerger_module.cc	C++ — rebuild	Pre-existing bug in base v04p04. PR pending.
NGMultiSliceCryoW hitInput	nugraph_icarus.fcl	FCL only — 3 copies	CryoW inherited CryoE loader inputs.
Pandora → cluster3DCryoE	stage1_icarus_defs. fcl	FCL only — 3 copies	Makes all hit sources consistent.
New FilterDecoder	larrecodnn (Ilena branch)	Checkout + rebuild	Fixes art::Ptr key → positional index mismatch.
lardataobj	lardataobj (Ilena branch)	Checkout + rebuild	Class definitions for NuGraph multi-slice data products.
2D model	nugraph_icarus.fcl	FCL only — 3 copies	Change modelFileName to model_mvmpmr_1201.pt. Add dir to FW_SEARCH_PATH.





APPENDIX A | NUGRAPH2 IN THE ICARUS/SBN RECONSTRUCTION FRAMEWORK



A.1 WHAT IS NUGRAPH2?

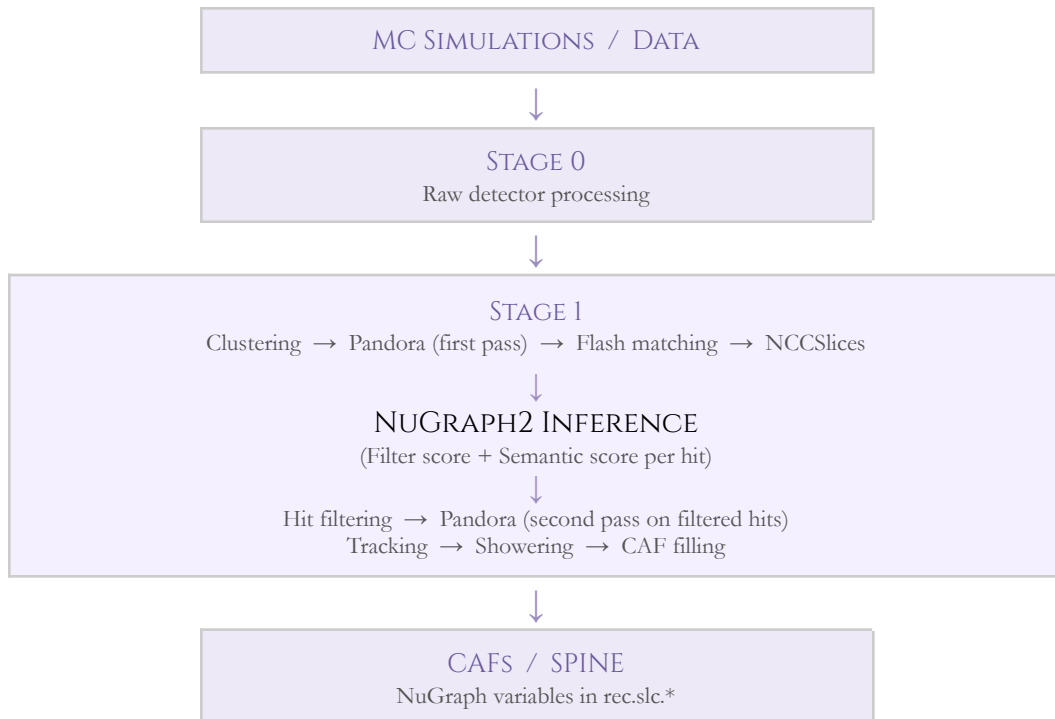
NuGraph2 is a Graph Neural Network (GNN) designed for low-level reconstruction of neutrino interactions in liquid argon time projection chamber (LArTPC) detectors. It operates at the individual TPC hit level, processing the raw geometry of hits within a Pandora slice and assigning each hit two scores:

- ≡ A filter score (binary: signal vs background) used to remove noise and cosmic hits before a second-pass Pandora reconstruction.
- ≡ A semantic score (multi-class): **MIP** (minimum ionising particle), **HIP** (highly ionising particle), **Shower**, **Michel**, **Diffuse**.

NuGraph2 enhances Pandora reconstruction by providing hit-level physics context that standard Pandora algorithms cannot access. By running GNN inference over the full hit graph of a neutrino candidate slice, it captures long-range geometric relationships between hits across all three wire planes simultaneously.

A.2 HOW NUGRAPH2 FITS INTO THE PIPELINE

NuGraph2 runs inside the Stage 1 reconstruction as an intermediate step between the first-pass Pandora reconstruction and the second-pass Pandora reconstruction on filtered hits. The full chain is:



A.3 WHY NUGRAPH2 IMPROVES RECONSTRUCTION EFFICIENCY

Standard Pandora slice reconstruction is limited to local geometric heuristics within each wire plane. NuGraph2 overcomes this by constructing a graph over all hits in a neutrino candidate slice, connecting hits across planes via shared SpacePoints, and running message-passing inference to propagate physics context globally.



- ≡ The filter decoder removes background and cosmic-ray hits before the second-pass Pandora sees them, reducing fragmentation and improving particle flow object (PFP) quality.
- ≡ The semantic decoder provides hit-level particle type classification (MIP/HIP/Shower/Michel/Diffuse) which can be used downstream for particle identification and selection.
- ≡ By running on neutrino candidate slices only (selected by NCCSlicesProducer after flash matching), NuGraph2 processes only the relevant portion of each event, keeping runtime manageable.

Note: NuGraph2 runs after Pandora slicing and flash matching, and before the second-pass Pandora reconstruction on filtered hits. It does not replace Pandora — it enhances it by providing hit-level physics scores that inform the second-pass reconstruction.

A.4 CAF VARIABLES PRODUCED

The following NuGraph2 variables are written to the CAF by sbncode:

CAF Variable	Description
<code>rec.slc.ng_filt_pass_frac</code>	Fraction of hits in slice passing the NuGraph filter score cut
<code>rec.slc.ng_plane[p].mip_hits</code>	Number of MIP-tagged hits per plane <i>p</i>
<code>rec.slc.ng_plane[p].hip_hits</code>	Number of HIP-tagged hits per plane <i>p</i>
<code>rec.slc.ng_plane[p].shr_hits</code>	Number of Shower-tagged hits per plane <i>p</i>
<code>rec.slc.ng_plane[p].mhl_hits</code>	Number of Michel-tagged hits per plane <i>p</i>
<code>rec.slc.ng_plane[p].dif_hits</code>	Number of Diffuse-tagged hits per plane <i>p</i>
<code>rec.slc.reco.pfp.ngscore.*</code>	Per-PFP NuGraph semantic scores



APPENDIX B | GREP CHEATSHEET

grep searches for patterns in files or command output. It is the fastest first diagnostic tool for any log file from a lar job.

B.1 BASIC SYNTAX

```
grep [OPTIONS] PATTERN FILE
grep [OPTIONS] PATTERN FILE1 FILE2 ...
command | grep [OPTIONS] PATTERN # pipe from another command
```

B.2 KEY OPTIONS

Option	Effect
-i	Case-insensitive match
-n	Print line numbers alongside matches
-v	Invert match -- print lines that do NOT match
-r	Recursive -- search all files in a directory tree
-l	Print only filenames of files containing a match (not the matching lines)
-c	Print only a count of matching lines per file
-A N	Print N lines After each match (context below)
-B N	Print N lines Before each match (context above)
-C N	Print N lines of Context around each match (both above and below)
-E	Extended regex -- enables , +, ?, () without escaping
-w	Match whole words only
-m N	Stop after N matches
--include=*.cc	Only search files matching a glob pattern (use with -r)

B.3 COMMON PATTERNS FOR LAR JOBS

Find errors and crashes in a log

```
grep -n 'error\|Error\|Abort\|Segmentation\|SIGSEGV\|Fatal\|MSG-e' job.log |
head -30
```

Track module execution with --trace

```
grep -n 'module for event\|finished for event' job.log | tail -30
```

Search across multiple files (3 FCL copies)

```
grep -n 'HitLabel' \
  srcs/icarusscode/icarusscode/TPC/NuGraph/nugraph_icarus.fcl \
  localProducts*/icarusscode/*/job/nugraph_icarus.fcl \
  build_slf7.x86_64/icarusscode/job/nugraph_icarus.fcl
```



Find a string recursively across all source files

```
grep -rn 'HitFinderModuleLabel' srcs/icaruscode/fcl/ | head -20
```

OR pattern -- multiple search terms

```
grep -n 'Invalid read\|Invalid write\|HitMerger\|FlashPredict' valgrind.log | head -40
```

Exclude lines matching a pattern

```
# Show all map lines that are not 7f... or ff... (memory map cleanup):
grep -v '^7f\|^ff' stage1.log | tail -50
```

Count hits before and after NuGraph filter

```
grep 'Number of hits' stage1_5events.log
```

Pipe grep into grep (AND logic)

```
# Lines containing both 'Invalid' and 'HitMerger':
grep 'Invalid' valgrind.log | grep 'HitMerger'
```

B.4 USEFUL ONE-LINERS

Show only the last N lines of grep output

```
grep 'pattern' file.log | tail -20
```

Show matching lines with filename (multi-file search)

```
grep -Hn 'combineHits' srcs/icaruscode/fcl/**/*.fcl
```

Extract a block of lines around a match

```
grep -n -A5 -B2 'ICARUSFilteredNuSliceHitsProducer' stage1.log | head -40
```

Check all three FCL copies for a value

```
for f in srcs/icaruscode/icaruscode/TPC/NuGraph/nugraph_icarus.fcl \
    localProducts*/icaruscode/*/job/nugraph_icarus.fcl \
    build_slf7.x86_64/icaruscode/job/nugraph_icarus.fcl; do
```



```
echo "=== $f ==="  
grep 'HitLabel' $f  
done
```



APPENDIX C | GDB CHEATSHEET

GDB (GNU Debugger) attaches to a running or crashed process and lets you inspect the call stack, variables, and memory at the moment of a crash. For lar jobs, it is used to find where a segfault occurred and what the call stack looked like.

C.1 SETUP FOR LAR JOBS

Always unset Python environment variables before running GDB inside the ICARUS container, or GDB's embedded Python will fail to load:

```
unset PYTHONPATH PYTHONHOME
```

C.2 RUNNING GDB

Interactive mode (run and inspect manually)

```
gdb --args lar -c my.fcl -n 1 -s input.root --process-name redo -o /dev/null
(gdb) run
(gdb) bt          # backtrace after crash
(gdb) bt full     # full backtrace with local variables
(gdb) quit
y                # confirm quit
```

Automated mode (log and exit)

```
gdb -ex "set logging file gdb_bt.log" \
    -ex "set logging on" \
    -ex "run" \
    -ex "bt" \
    -ex "bt full" \
    -ex "quit" \
    --args lar -c my.fcl -n 1 -s input.root --process-name redo -o /dev/null
2>&1
```

C.3 KEY GDB COMMANDS

Command	Effect
run	Start execution of the program
bt	Print backtrace (call stack at crash point)
bt full	Print backtrace with local variables for each frame
frame N	Switch to stack frame N
info locals	Print all local variables in the current frame
print VAR	Print the value of a variable
print *ptr	Dereference and print the value at a pointer
break FUNC	Set a breakpoint at the start of a function
break FILE:LINE	Set a breakpoint at a specific file and line number
continue	Continue execution after a breakpoint
next	Execute the next line (step over function calls)



Command	Effect
step	Execute the next line (step into function calls)
info args	Print function arguments in the current frame
list	Show source code around the current line
quit	Exit GDB (confirm with y if a process is active)
set logging file F	Redirect all GDB output to file F
set logging on	Enable logging to the file set above

C.4 READING A BACKTRACE

A backtrace shows the chain of function calls leading to the crash, starting from frame #0 (the innermost call, where the crash occurred) upwards. Example:

```
#0 _int_malloc ()                <- crash here: heap allocator corrupted
#1 malloc ()
#2 operator new ()
#3 art::detail::IPRHelper<...>    <- FindManyP construction
#4 ICARUSFilteredNuSliceHitsProducer::produce
#5 art::EDProducer::doEvent
...
```

In heap corruption crashes, the crash site (#0) is almost always inside malloc or _int_malloc. The actual bug (the bad write) happened earlier, in a higher-numbered frame. Look for the first frame that references your own code or art association code.

C.5 COMMON PATTERNS

Heap corruption (the crash site != the bug site)

```
#0 _int_malloc      <- crash site (heap bookkeeping corrupted)
#1 malloc
#N YourModule::produce <- the actual bad write happened somewhere upstream
```

Heap corruption: use valgrind to find the actual write site. GDB only shows where the corrupted heap was first accessed, not where it was corrupted.

Null pointer dereference

```
#0 SomeFunction (this=0x0 ...)    <- this=0x0 means null pointer
```

Out-of-bounds vector access

```
#0 std::vector<...>::operator[]    <- usually shows size vs index
    index = 4532, size = 200      <- visible with bt full or print
```

Variables show as <optimised out>

This happens in prof (optimised) builds. Rebuild with -q e26:debug using mrb newDev -T debug_build -q e26:debug -f to get full variable visibility.



C.6 SETTING BREAKPOINTS FOR NUGRAPH DEBUGGING

```
# Break at the start of FilterDecoder::writeToEvent:
(gdb) break FilterDecoder::writeToEvent
(gdb) run

# Once stopped, inspect idsmmap size and max key:
(gdb) info locals
(gdb) print idsmmap.size()
(gdb) print size

# Continue to next breakpoint or crash:
(gdb) continue
```